

---

---

# Seq2seq, Attention, Self attention, Transformer, BERT

---

---

Tuan Nguyen - AI4E

---

---



# Outline

- RNN review
- Seq2seq
- Beam search
- } Attention
- } Self-attention
- } Transformer
- { BERT


# Recurrent Neural Network

Sequence  
time - semi .

Usually drawn as:



# RNN Formula




The state consists of a single “*hidden*” vector  $\mathbf{h}$ :

$$h_t = f_{\underline{W}}(h_{t-1}, x_t)$$



$$\left. \begin{array}{l} h_t = \tanh(\underline{W}_{hh} h_{t-1} + \underline{W}_{xh} x_t) \\ y_t = \underline{W}_{hy} h_t \end{array} \right\}$$

# Forward





$$h_t = \tanh W \begin{pmatrix} x_t \\ h_{t-1} \end{pmatrix}$$

$$y_t = F(h_t)$$


$$C_t = \text{Loss}(y_t, \text{GT}_t)$$

# Deep RNN



# Recurrent neural network problem

one to one



NN  
CNN

one to many






Image.  
capturing

many to one




video classif.

many to many



translati.

many to many




charact. generatn.

# Character-level language model example

Vocabulary:  
[h,e,l,o]


Example training  
sequence:  
**“hello”**



# Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training  
sequence:  
**“hello”**




# Character-level language model example

Vocabulary:  
[h,e,l,o]

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$


Example training  
sequence:  
**“hello”**




# Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training  
sequence:  
**“hello”**



# Long short term memory (LSTM) / GRU .



# Translation

Dịch tiếng Anh  
Translate English-Vietnamese 

English Translate   

I would like to improve my English **skills**

Dictionary  [Basic] [Technical]  On/Off

 skills  
skill /skil/  
• *danh từ*  
◦ **sự khéo léo, sự khéo tay, sự tinh xảo; kỹ năng, kỹ xảo**

Tiếng Việt


Tôi muốn cải thiện kỹ năng tiếng anh của mình

Example  On/Off

 Next >>


- I'm expecting to gain more **skills** and knowledge, Sir.
- I don't have any **skills** at fishing.
- I humbly offer my dedication to practice the medical **skills** and knowledge I gained from 10 long years of study.

# Model




# Seq2seq



autoreular.




# Seq2Seq



# Seq2seq



# Seq2seq - prediction



# Greedy search


	Timestep	1	2	3	4
A	0.5	0.1	0.2	0.0	
B	0.2	0.4	0.2	0.2	
C	0.2	0.3	0.4	0.2	
<eos>	0.1	0.2	0.2	0.6	

Beam search,  $k=2$ , hyper-parameter  $\gamma$  và  $k$  nhanh tốt nhất


$\Sigma^n$ ,  $n$  là length output.

$\sqrt[5]{y_1 \cdot y_2 \cdot y_3 \cdot y_4 \cdot y_5}$

$\sqrt[3]{y_1 \cdot y_2 \cdot y_3}$





# Word embedding (glove, word2vec)




# Word semantic


$$\text{que} - \text{king} + \text{man} \approx \text{woman}$$



Male-Female




Verb tense





Country-Capital

# Attention - motivation

The good ~~b~~ ~~god~~ → P<sub>IS</sub> / neg



# Attention



Handwritten notes calculating the weighted sum of values:


$$V = \sum v_i \cdot a_i$$

where  $a_i$  are the attention weights and  $v_i$  are the values from the Memory.

$a_i = \alpha(q, k_i)$  : độ quan trọng của từ i với từ cần tìm,

$$b_i = \frac{\exp(a_i)}{\sum_j \exp(a_j)}, \mathbf{b} = [b_1, \dots, b_n]^T$$

# Attention





# Attention function


$$a_i = \alpha(q_i k_i) \cdot \alpha: q_i k_i \rightarrow a_i$$


Name	Alignment score function
Content-base attention	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \text{cosine}[\mathbf{s}_t, \mathbf{h}_i]$
Additive(*)	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \mathbf{v}_a^\top \tanh(\mathbf{W}_a[\mathbf{s}_t; \mathbf{h}_i])$
Location-Base	$\alpha_{t,i} = \text{softmax}(\mathbf{W}_a \mathbf{s}_t)$ Note: This simplifies the softmax alignment to only depend on the target position.
General	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \mathbf{s}_t^\top \mathbf{W}_a \mathbf{h}_i$ where $\mathbf{W}_a$ is a trainable weight matrix in the attention layer.
Dot-Product	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \mathbf{s}_t^\top \mathbf{h}_i$
Scaled Dot-Product(^)	$\text{score}(\mathbf{s}_t, \mathbf{h}_i) = \frac{\mathbf{s}_t^\top \mathbf{h}_i}{\sqrt{n}}$ Note: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.


**Encoder**












# Self attention

: hìn. high-level feature cuộn.

Self-attention



input #1

1	0	1	0
---	---	---	---





input #2

0	2	0	2
---	---	---	---

input #3

1	1	1	1
---	---	---	---

Self-attention




I

love

you

Self-attention



input #1

1	0	1	0
---	---	---	---


input #2

0	2	0	2
---	---	---	---


input #3

1	1	1	1
---	---	---	---


Self-attention




### Self-attention




## Self-attention




## Self-attention




## Self-attention



## Self-attention







# Transformer




# Seq2seq




# Transformer







ENCODER #2





ENCODER #1

Feed Forward  
Neural Network

Feed Forward  
Neural Network





Self-Attention




Thinking

Machines









Decoding time step: 1 2 3 4 5 6


OUTPUT



EMBEDDING  
WITH TIME  
SIGNAL



EMBEDDINGS




INPUT      Je      suis      étudiant


Decoding time step: 1 2 3 4 5 6

OUTPUT

|





# Transformer




# Multi-head attention

- 1) This is our input sentence\*
- 2) We embed each word\*
- 3) Split into 8 heads. We multiply  $X$  or  $R$  with weight matrices
- 4) Calculate attention using the resulting  $Q/K/V$  matrices
- 5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer




\* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one




# BERT

## GLUE scores evolution over 2018-2019




# BERT model




- Mask Language Modeling (MLM)
- Next Sentence Prediction (NSP)

# Masked LM (MLM)



# Next sentence prediction



# Fine-tuning BERT

- Classification tasks such as sentiment analysis.
- In Question Answering tasks (e.g. SQuAD v1.1).
- In Named Entity Recognition (NER).

# Q&A



@LateNightSeth

